

Problem Statement :

Customer churn or customer attrition is a tendency of clients or customers to abandon a brand and stop being a paying client of a particular business or organization. The percentage of customers that discontinue using a company's services or products during a specific period is called a customer churn rate. Several bad experiences (or just one) are enough, and a customer may quit. And if a large chunk of unsatisfied customers churn at a time interval, both material losses and damage to reputation would be enormous.

A reputed bank "ABC BANK" wants to predict the Churn rate. Create a model by using different machine learning approaches that can predict the best result.

Dataset Description :

This is a public dataset, The dataset format is given below.

Inside the dataset, there are 10000 rows and 14 different columns.

The target column here is Exited here.

```
import pandas as pd

df=pd.read_csv("https://raw.githubusercontent.com/sagnikghoshcr7/Bank-
Customer-Churn-Prediction/master/Churn_Modelling.csv")

df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	2	0.00	1	1	1
1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1

```
3      93826.63      0
4      79084.10      0
```

```
df.tail()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
9995	9996	15606229	Obijiaku	771	France	Male
39						
9996	9997	15569892	Johnstone	516	France	Male
35						
9997	9998	15584532	Liu	709	France	Female
36						
9998	9999	15682355	Sabbatini	772	Germany	Male
42						
9999	10000	15628319	Walker	792	France	Female
28						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
9995	5	0.00	2	1		0
9996	10	57369.61	1	1		1
9997	7	0.00	1	0		1
9998	3	75075.31	2	1		0
9999	4	130142.79	1	1		0

	EstimatedSalary	Exited
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
      'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
      'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
#Dropping Irrelevant Features
```

```
df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], inplace=True)
```

```
-----
-----
KeyError                                Traceback (most recent call
last)
<ipython-input-11-075cb218a943> in <cell line: 2>()
      1 #Dropping Irrelevant Features
----> 2
```

```

df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], inplace=True)

/usr/local/lib/python3.10/dist-packages/pandas/util/_decorators.py in
wrapper(*args, **kwargs)
    329         stacklevel=find_stack_level(),
    330     )
--> 331     return func(*args, **kwargs)
    332
    333     # error: "Callable[[VarArg(Any), KwArg(Any)], Any]"
has no

/usr/local/lib/python3.10/dist-packages/pandas/core/frame.py in
drop(self, labels, axis, index, columns, level, inplace, errors)
    5397         weight 1.0      0.8
    5398         """
-> 5399     return super().drop(
    5400         labels=labels,
    5401         axis=axis,

/usr/local/lib/python3.10/dist-packages/pandas/util/_decorators.py in
wrapper(*args, **kwargs)
    329         stacklevel=find_stack_level(),
    330     )
--> 331     return func(*args, **kwargs)
    332
    333     # error: "Callable[[VarArg(Any), KwArg(Any)], Any]"
has no

/usr/local/lib/python3.10/dist-packages/pandas/core/generic.py in
drop(self, labels, axis, index, columns, level, inplace, errors)
    4503     for axis, labels in axes.items():
    4504         if labels is not None:
-> 4505         obj = obj._drop_axis(labels, axis,
level=level, errors=errors)
    4506
    4507         if inplace:

/usr/local/lib/python3.10/dist-packages/pandas/core/generic.py in
_drop_axis(self, labels, axis, level, errors, only_slice)
    4544         new_axis = axis.drop(labels, level=level,
errors=errors)
    4545     else:
-> 4546         new_axis = axis.drop(labels, errors=errors)
    4547         indexer = axis.get_indexer(new_axis)
    4548

/usr/local/lib/python3.10/dist-packages/pandas/core/indexes/base.py in
drop(self, labels, errors)
    6932     if mask.any():
    6933         if errors != "ignore":

```

```
-> 6934                 raise KeyError(f"{list(labels[mask])} not
found in axis")
    6935                 indexer = indexer[~mask]
    6936                 return self.delete(indexer)
```

```
KeyError: "['RowNumber', 'CustomerId', 'Surname'] not found in axis"
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance
0	619	France	Female	42	2	0.00
1						
1	608	Spain	Female	41	1	83807.86
1						
2	502	France	Female	42	8	159660.80
3						
3	699	France	Female	39	1	0.00
2						
4	850	Spain	Female	43	2	125510.82
1						

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	CreditScore	10000 non-null	int64
1	Geography	10000 non-null	object
2	Gender	10000 non-null	object
3	Age	10000 non-null	int64
4	Tenure	10000 non-null	int64
5	Balance	10000 non-null	float64
6	NumOfProducts	10000 non-null	int64
7	HasCrCard	10000 non-null	int64
8	IsActiveMember	10000 non-null	int64
9	EstimatedSalary	10000 non-null	float64
10	Exited	10000 non-null	int64

```
dtypes: float64(2), int64(7), object(2)
```

```
memory usage: 859.5+ KB
```

```
df.describe()
```

	CreditScore	Age	Tenure	Balance
NumOfProducts \				
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288
std	96.653299	10.487806	2.892174	62397.405202
min	350.000000	18.000000	0.000000	0.000000
25%	584.000000	32.000000	3.000000	0.000000
50%	652.000000	37.000000	5.000000	97198.540000
75%	718.000000	44.000000	7.000000	127644.240000
max	850.000000	92.000000	10.000000	250898.090000

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.70550	0.515100	100090.239881	0.203700
std	0.45584	0.499797	57510.492818	0.402769
min	0.000000	0.000000	11.580000	0.000000
25%	0.000000	0.000000	51002.110000	0.000000
50%	1.000000	1.000000	100193.915000	0.000000
75%	1.000000	1.000000	149388.247500	0.000000
max	1.000000	1.000000	199992.480000	1.000000

```
#value counts()
df["Geography"].value_counts()
```

```
France      5014
Germany     2509
Spain       2477
Name: Geography, dtype: int64
```

```
#value counts()
df["Gender"].value_counts()
```

```
Male        5457
Female      4543
Name: Gender, dtype: int64
```

Encoding Categorical Data

```
df=pd.get_dummies(df,columns=['Geography','Gender'],drop_first=True)
df.head()
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	

	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	\
0	1	101348.88	1	0	
1	1	112542.58	0	0	
2	0	113931.57	1	0	
3	0	93826.63	0	0	
4	1	79084.10	0	0	

	Geography_Spain	Gender_Male
0	0	0
1	1	0
2	0	0
3	0	0
4	1	0

Some insights about the target variable

```
df["Exited"].value_counts()
```

```
0    7963
```

```
1    2037
```

```
Name: Exited, dtype: int64
```

```
X=df.drop(['Exited'],axis=1)
```

```
y=df["Exited"]
```

```
X
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	

```
IsActiveMember EstimatedSalary Geography_Germany
```

Geography_Spain \			
0	1	101348.88	0
0			
1	1	112542.58	0
1			
2	0	113931.57	0
0			
3	0	93826.63	0
0			
4	1	79084.10	0
1			
...
...			
9995	0	96270.64	0
0			
9996	1	101699.77	0
0			
9997	1	42085.58	0
0			
9998	0	92888.52	1
0			
9999	0	38190.78	0
0			

Gender_Male	
0	0
1	0
2	0
3	0
4	0
...	...
9995	1
9996	1
9997	0
9998	1
9999	0

[10000 rows x 11 columns]

y

0	1
1	0
2	1
3	0
4	0
...	..
9995	0
9996	0
9997	1

```
9998      1
9999      0
Name: Exited, Length: 10000, dtype: int64
```

Keras:

Keras is a high-level, deep learning API developed by Google for implementing neural networks. It is written in Python and is used to make the implementation of neural networks easy. It also supports multiple backend neural network computation.

Keras allows you to switch between different back ends. The frameworks supported by Keras are:

Tensorflow

Theano

PlaidML

MXNet

CNTK (Microsoft Cognitive Toolkit)

Why Do We Need Keras?

Keras is an API that was made to be easy to learn for people. Keras was made to be simple. It offers consistent & simple APIs, reduces the actions required to implement common code, and explains user error clearly.

Prototyping time in Keras is less. This means that your ideas can be implemented and deployed in a shorter time. Keras also provides a variety of deployment options depending on user needs.

Languages with a high level of abstraction and inbuilt features are slow and building custom features in then can be hard. But Keras runs on top of TensorFlow and is relatively fast. Keras is also deeply integrated with TensorFlow, so you can create customized workflows with ease.

The research community for Keras is vast and highly developed. The documentation and help available are far more extensive than other deep learning frameworks.

Keras is used commercially by many companies like Netflix, Uber, Square, Yelp, etc which have deployed products in the public domain which are built using Keras.

Apart from this, Keras has features such as :

It runs smoothly on both CPU and GPU.

It supports almost all neural network models.

It is modular in nature, which makes it expressive, flexible, and apt for innovative research.

Applications of Keras

Keras is used for creating deep models which can be productized on smartphones.

Keras is also used for distributed training of deep learning models.

Keras is used by companies such as Netflix, Yelp, Uber, etc.

Keras is also extensively used in deep learning competitions to create and deploy working models, which are fast in a short amount of time.

```
import tensorflow
from tensorflow import keras

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense

len(X.columns)

11

model=Sequential()
model.add(Dense(11,activation="sigmoid",input_dim=11))#input layer
model.add(Dense(11,activation="sigmoid"))#hidden layer
model.add(Dense(1,activation="sigmoid"))#output layer
```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 11)	132
dense_1 (Dense)	(None, 11)	132
dense_2 (Dense)	(None, 1)	12
Total params: 276		
Trainable params: 276		
Non-trainable params: 0		

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
```

```

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()

X_train_scale=scaler.fit_transform(X_train)
X_test_scale=scaler.transform(X_test)

#The model.compile function in Keras is used to specify the optimizer,
loss function, and evaluation metrics that will be used during the
training of the neural network model.

model.compile(optimizer="Adam",
              loss="binary_crossentropy",
              metrics=["accuracy"]

)

model.fit(X_train_scale,y_train,batch_size=50,epochs=10,verbose=1)

Epoch 1/10
160/160 [=====] - 1s 2ms/step - loss: 0.4325
- accuracy: 0.8074
Epoch 2/10
160/160 [=====] - 0s 2ms/step - loss: 0.4297
- accuracy: 0.8102
Epoch 3/10
160/160 [=====] - 0s 2ms/step - loss: 0.4279
- accuracy: 0.8105
Epoch 4/10
160/160 [=====] - 0s 2ms/step - loss: 0.4263
- accuracy: 0.8114
Epoch 5/10
160/160 [=====] - 0s 2ms/step - loss: 0.4250
- accuracy: 0.8119
Epoch 6/10
160/160 [=====] - 0s 2ms/step - loss: 0.4235
- accuracy: 0.8155
Epoch 7/10
160/160 [=====] - 0s 2ms/step - loss: 0.4220
- accuracy: 0.8174
Epoch 8/10
160/160 [=====] - 0s 2ms/step - loss: 0.4206
- accuracy: 0.8191
Epoch 9/10
160/160 [=====] - 0s 2ms/step - loss: 0.4193
- accuracy: 0.8207
Epoch 10/10
160/160 [=====] - 0s 2ms/step - loss: 0.4176
- accuracy: 0.8234

<keras.callbacks.History at 0x78b4377885b0>

```

X_train_scale: Your scaled training data.

y_train: The corresponding target labels for your training data.

batch_size: The number of samples to use in each iteration of gradient descent during training. In this case, you're using a batch size of 50, which means that the model will update its weights after every 50 samples.

epochs: The number of times the training dataset will be passed through the model during training. You're running 10 epochs, so the model will go through the entire training dataset 10 times.

verbose: This controls the verbosity of the training process. Setting it to 1 means that progress bars and logs will be displayed during training, providing you with feedback on the training progress.

```
y_pred=model.predict(X_test_scale)
63/63 [=====] - 0s 2ms/step
y_pred
array([[0.19777617],
       [0.36974898],
       [0.14303005],
       ...,
       [0.1319361 ],
       [0.09910323],
       [0.14592502]], dtype=float32)
y_pred=y_pred.argmax(axis=-1)
from sklearn.metrics import accuracy_score #used to calculate the
accuracy of classification models
accuracy_score(y_test,y_pred)
0.7975
```